

A C++ MQTT Message Broker for the Enterprise

John Dubchak
Castlight Health Inc.
Copyright © 2016

Agenda

- Shameless self-promotion
- Overview of Messaging
 - Basic messaging concepts
 - Enterprise Messaging
- The MQTT Messaging Protocol
- Designing and Implementing an MQTT Broker

ABOUT ME

The Basics

MESSAGING OVERVIEW

What is Messaging?

The communication of data from sender to receiver over a given messaging channel in an agreed upon format of exchange

What is Messaging?

The **communication** of data from sender to receiver over a given messaging channel in an agreed upon format of exchange

What is Messaging?

The communication of **data** from sender to receiver over a given messaging channel in an agreed upon format of exchange

What is Messaging?

The communication of data from **sender** to receiver over a given messaging channel in an agreed upon format of exchange

What is Messaging?

The communication of data from sender to **receiver** over a given messaging channel in an agreed upon format of exchange

What is Messaging?

The communication of data from sender to receiver over a given **messaging channel** in an agreed upon format of exchange

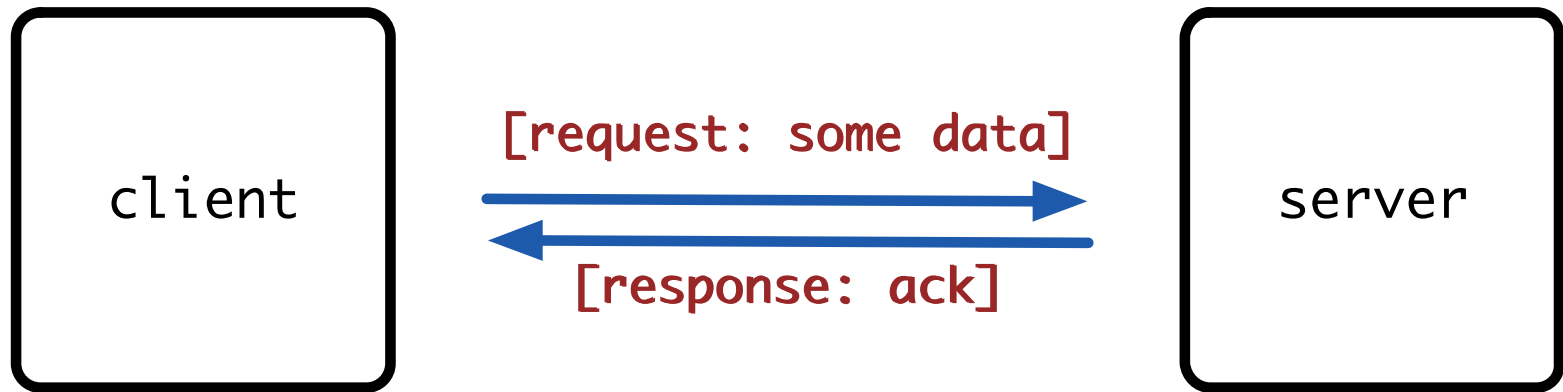
What is Messaging?

The communication of data from sender to receiver over a given messaging channel in an **agreed upon** format of exchange

What is Messaging?

The communication of data from sender to receiver over a given messaging channel in an agreed upon format of **exchange**

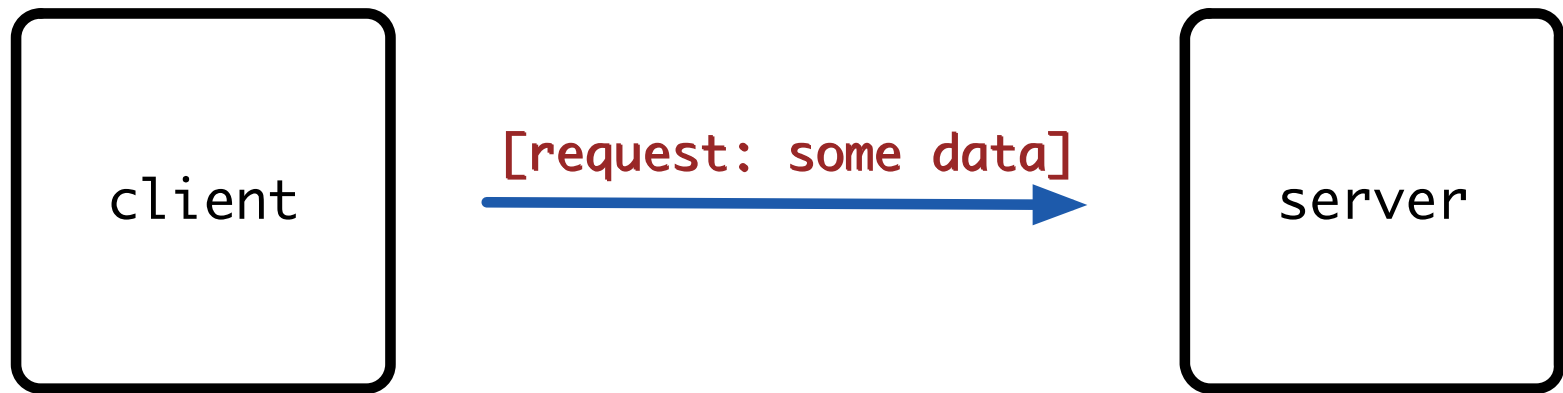
The Canonical Messaging Model




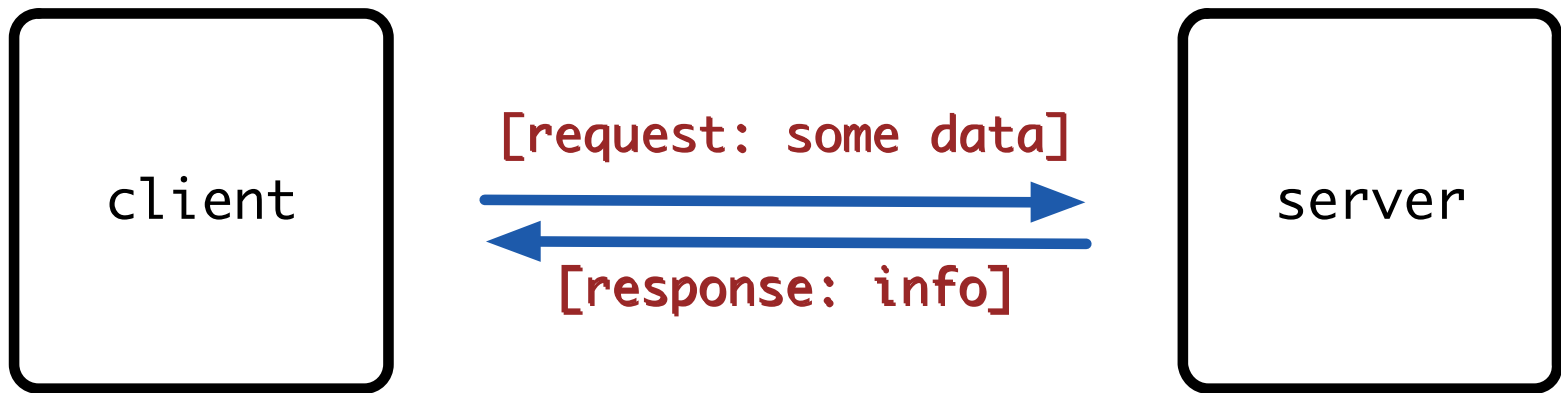
Messaging Architecture

STYLES OF COMMUNICATION

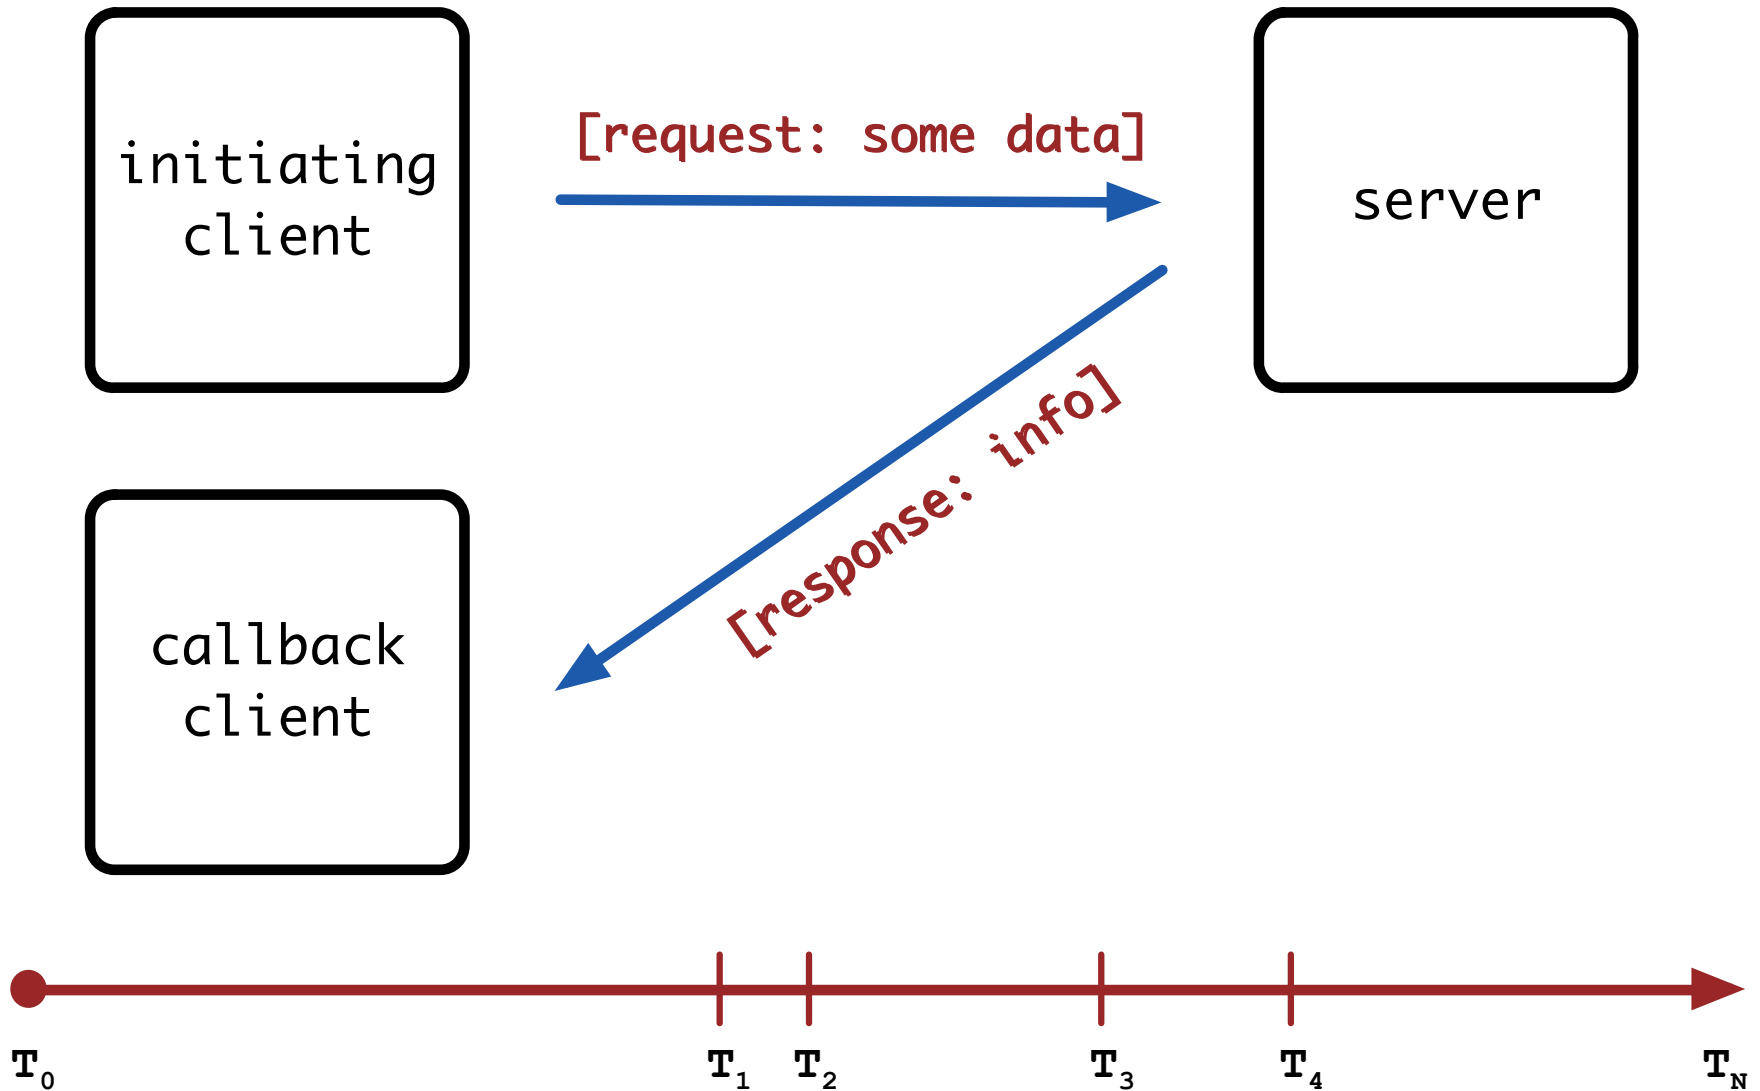
Point-to-Point Messaging



Request-Response



Request-Callback

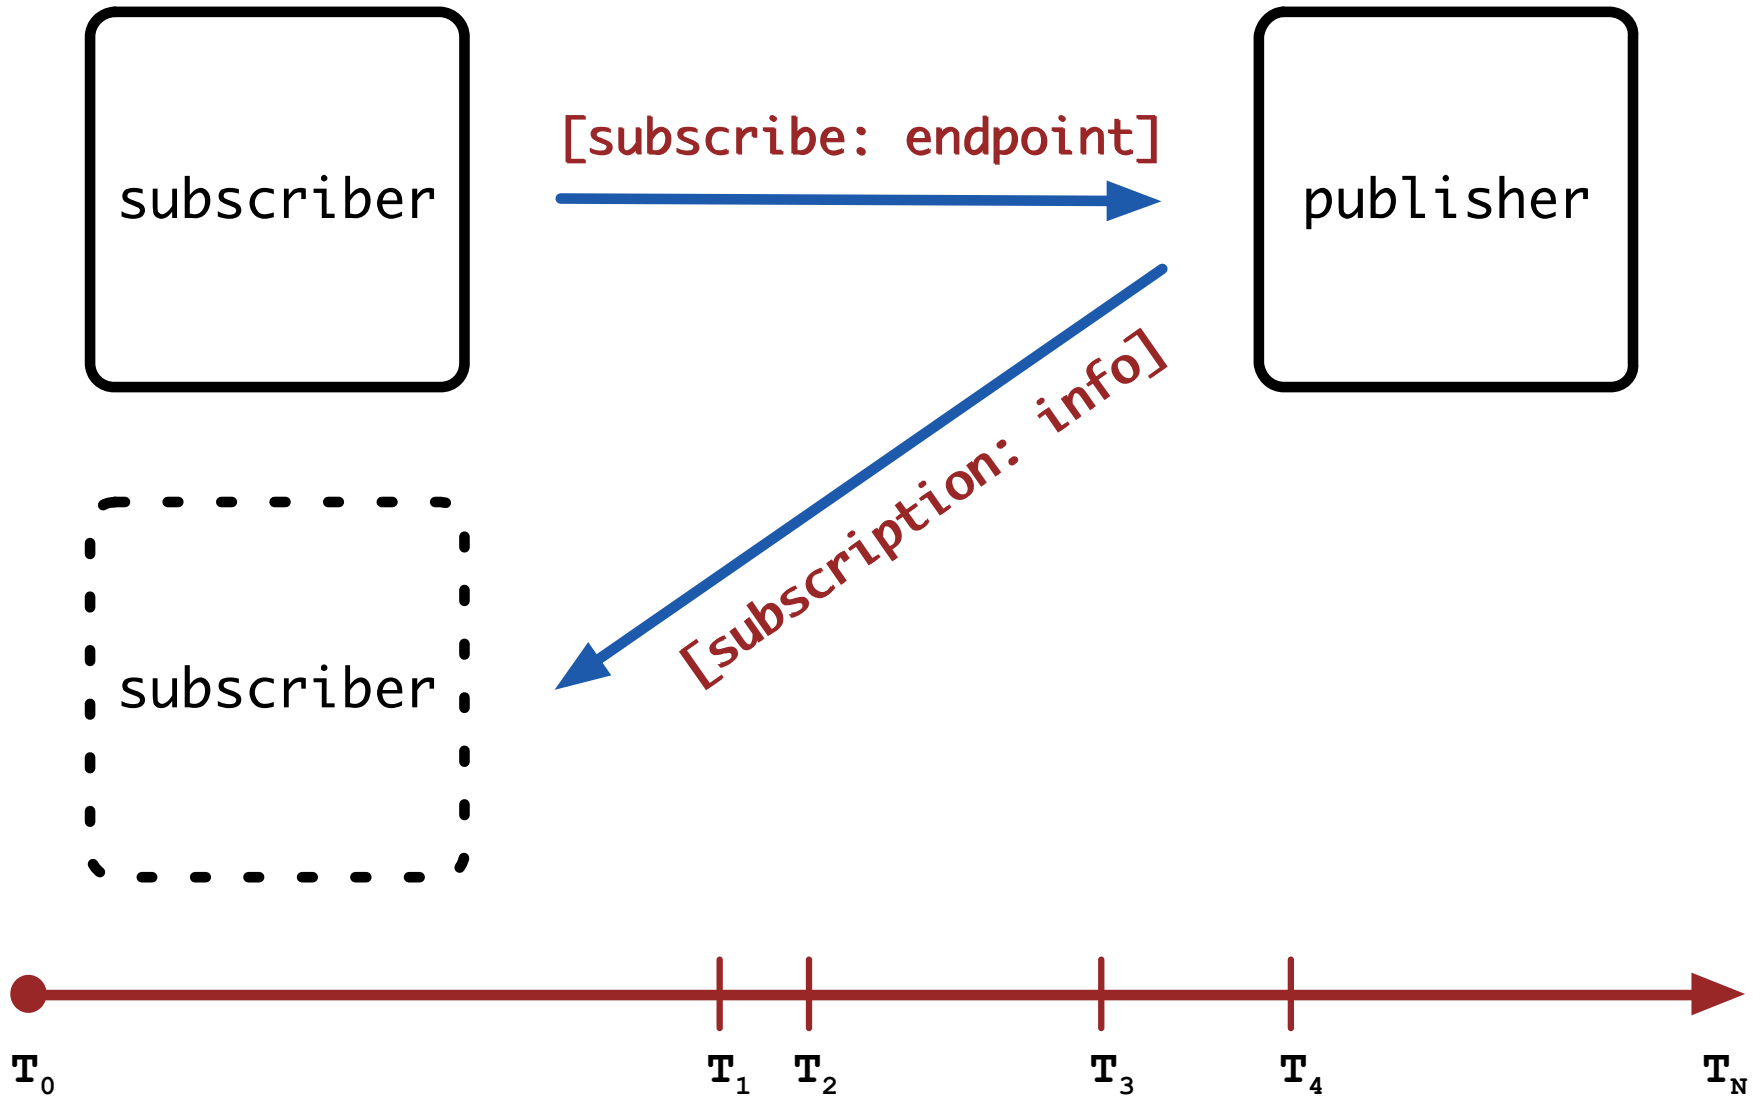


Actor Model

“The actor model in computer science is a mathematical model of concurrent computation that treats **actors** as the universal primitives of concurrent computation.”

https://en.wikipedia.org/wiki/Actor_model

Publish-Subscribe



ENTERPRISE MESSAGING

What is Enterprise Messaging?

- A suite of tools providing:
 - Business Process Orchestration
 - Systems and Data Integration
 - Monitoring
 - Transformation/Routing
 - Logging
- You may see MSMQ, IBM MQ Series, JMS Technologies
- Systems are either heterogeneous or homogenous

Architectural Attributes

- Highly Available
- Fault tolerant
- Secure
- Redundant
- Reliable/Delivery Retry
- Guaranteed Message Delivery
- Message Ordering
- Client Session Management

The Messaging Layer

THE MQTT PROTOCOL

What is MQTT?

- Message Queue Telemetry Transport
- Pub/Sub Message Protocol
- Standardized in OASIS & ISO
- Lightweight
- Constrained Vocabulary

Protocol Requirements


- Integer data values are 16-bits in Big-Endian order
- All text is UTF-8 encoded strings
- All strings are prefixed with a two byte length
- Strings are limited to 65,535 bytes
- Character data must be well-formed
- The NULL character '\0' is not permitted in String data

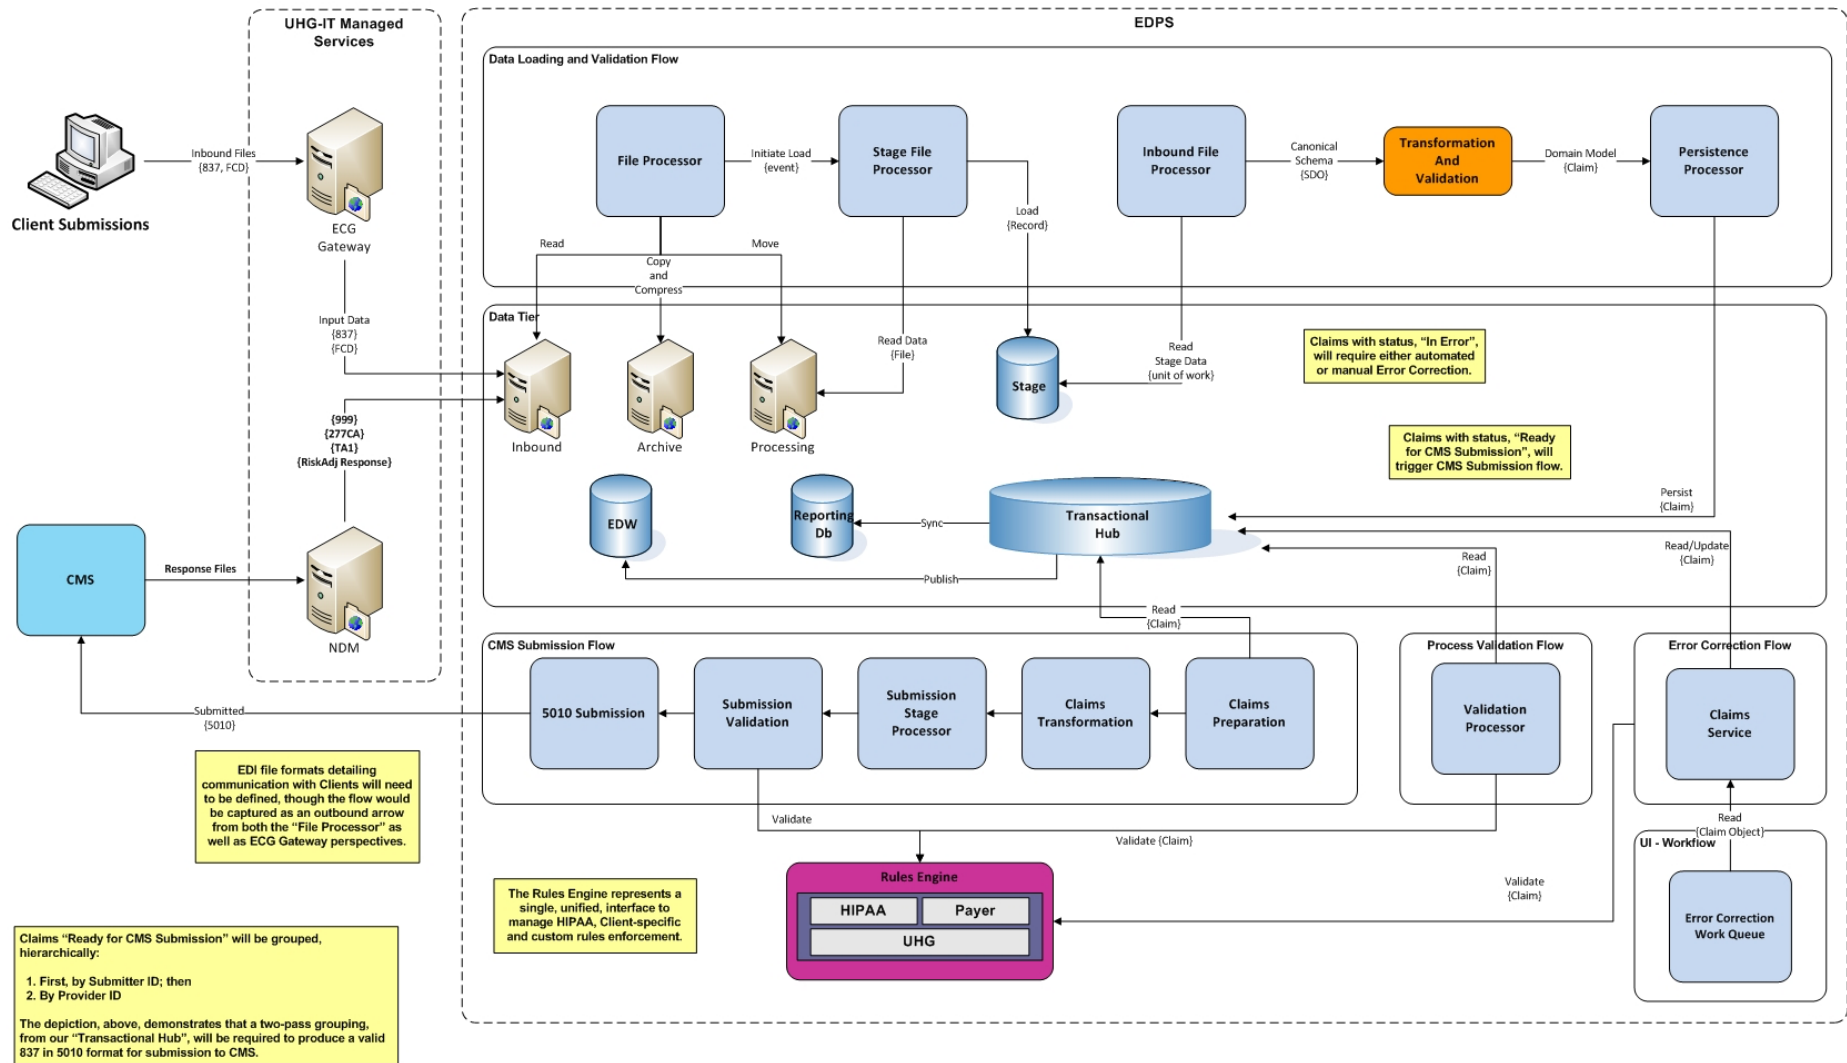
Message Types

- Connect
- Subscribe
- Unsubscribe
- Publish
- PubRel
- PubComp
- PingReq
- Disconnect
- ConnAck
- Suback
- Unsuback
- Puback
- PingResp

MQTT MESSAGING USE- CASES

Internet of Things







Making a connection

MQTT CLIENT CONNECT TO BROKER

MQTT Connect



MQTT Client Connect



PODs Used for Protocol Connect Data Communications

MQTT ABSTRACTIONS

```
struct Connect {

    std::uint8_t type      : 4;
    std::uint8_t           : 4;
    std::uint8_t length    : 8;

    char* protocol;
    char* level;

    union {
        unsigned char all;
        struct {
            std::uint8_t           : 1;
            std::uint8_t cleanSession : 1;
            std::uint8_t will      : 1;
            std::uint8_t qos       : 2;
            std::uint8_t will_retain : 1;
            std::uint8_t password  : 1;
            std::uint8_t username  : 1;
        } flags;
    } bits;

    char* clientId;
    char* topic;
    char* message;
    char* username;
    char* password;
};
```

```
struct Connect {
```

```
    std::uint8_t type      : 4;  
    std::uint8_t          : 4;  
    std::uint8_t length    : 8;
```

```
    char* protocol;
```

```
    char* level;
```

```
    union {
```

```
        unsigned char all;
```

```
        struct {
```

```
            std::uint8_t          : 1;
```

```
            std::uint8_t cleanSession : 1;
```

```
            std::uint8_t will        : 1;
```

```
            std::uint8_t qos          : 2;
```

```
            std::uint8_t will_retain  : 1;
```

```
            std::uint8_t password     : 1;
```

```
            std::uint8_t username     : 1;
```

```
        } flags;
```

```
    } bits;
```

```
    char* clientId;
```

```
    char* topic;
```

```
    char* message;
```

```
    char* username;
```

```
    char* password;
```

```
};
```

```
struct Connect {
```

```
    std::uint8_t type      : 4;  
    std::uint8_t          : 4;  
    std::uint8_t length    : 8;
```

```
    char* protocol;
```

```
    char* level;
```

```
    union {
```

```
        unsigned char all;
```

```
        struct {
```

```
            std::uint8_t          : 1;
```

```
            std::uint8_t cleanSession : 1;
```

```
            std::uint8_t will        : 1;
```

```
            std::uint8_t qos          : 2;
```

```
            std::uint8_t will_retain  : 1;
```

```
            std::uint8_t password     : 1;
```

```
            std::uint8_t username     : 1;
```

```
        } flags;
```

```
    } bits;
```

```
    char* clientId;
```

```
    char* topic;
```

```
    char* message;
```

```
    char* username;
```

```
    char* password;
```

```
};
```

```

struct Connect {

    std::uint8_t type      : 4;
    std::uint8_t          : 4;
    std::uint8_t length    : 8;

    char* protocol;
    char* level;

    union {
        unsigned char all;
        struct {
            std::uint8_t          : 1;
            std::uint8_t cleanSession : 1;
            std::uint8_t will      : 1;
            std::uint8_t qos       : 2;
            std::uint8_t will_retain : 1;
            std::uint8_t password  : 1;
            std::uint8_t username  : 1;
        } flags;
    } bits;

    char* clientId;
    char* topic;
    char* message;
    char* username;
    char* password;
};

```

```

struct Connect {

    std::uint8_t type      : 4;
    std::uint8_t          : 4;
    std::uint8_t length    : 8;

    char* protocol;
    char* level;

    union {
        unsigned char all;
        struct {
            std::uint8_t          : 1;
            std::uint8_t cleanSession : 1;
            std::uint8_t will      : 1;
            std::uint8_t qos       : 2;
            std::uint8_t will_retain : 1;
            std::uint8_t password  : 1;
            std::uint8_t username  : 1;
        } flags;
    } bits;

    char* clientId;
    char* topic;
    char* message;
    char* username;
    char* password;
};

```

Connection Flags

Flag Name	Purpose
Clean Session	0: MUST resume communication based on current session of Client Identifier
Will	1: If connection is accepted a message must be stored that is delivered upon subsequent disconnection 0: No message is to be stored
Will QoS	If Will Message is 0, Will QoS MUST be 0. If Will Message is 1, Will QoS can be either 0 (0x00), 1 (0x01), or 2 (0x02)
Will Retain	Broker should retain Will Message or not
Password	0, password must not be present, 1 it is
User Name	0, user name must not be present, 1 it is

```
struct ConnAck {  
    std::uint8_t type          :  
4;  
    std::uint8_t               :  
4;  
    std::uint8_t length        :  
8;  
    std::uint8_t               :  
7;  
    std::uint8_t session       :  
1;  
    std::uint8_t return_code   :  
8;  
};
```

* Session and return code are only used in ACK


```

struct ConnAck {
    std::uint8_t type          :
4;
    std::uint8_t              :
4;
    std::uint8_t length        :
8;
    std::uint8_t              :
7;
    std::uint8_t session       :
1;
    std::uint8_t return_code   :
8;
};

```

* Session and return code are only used in ACK

CONNACK Return Codes

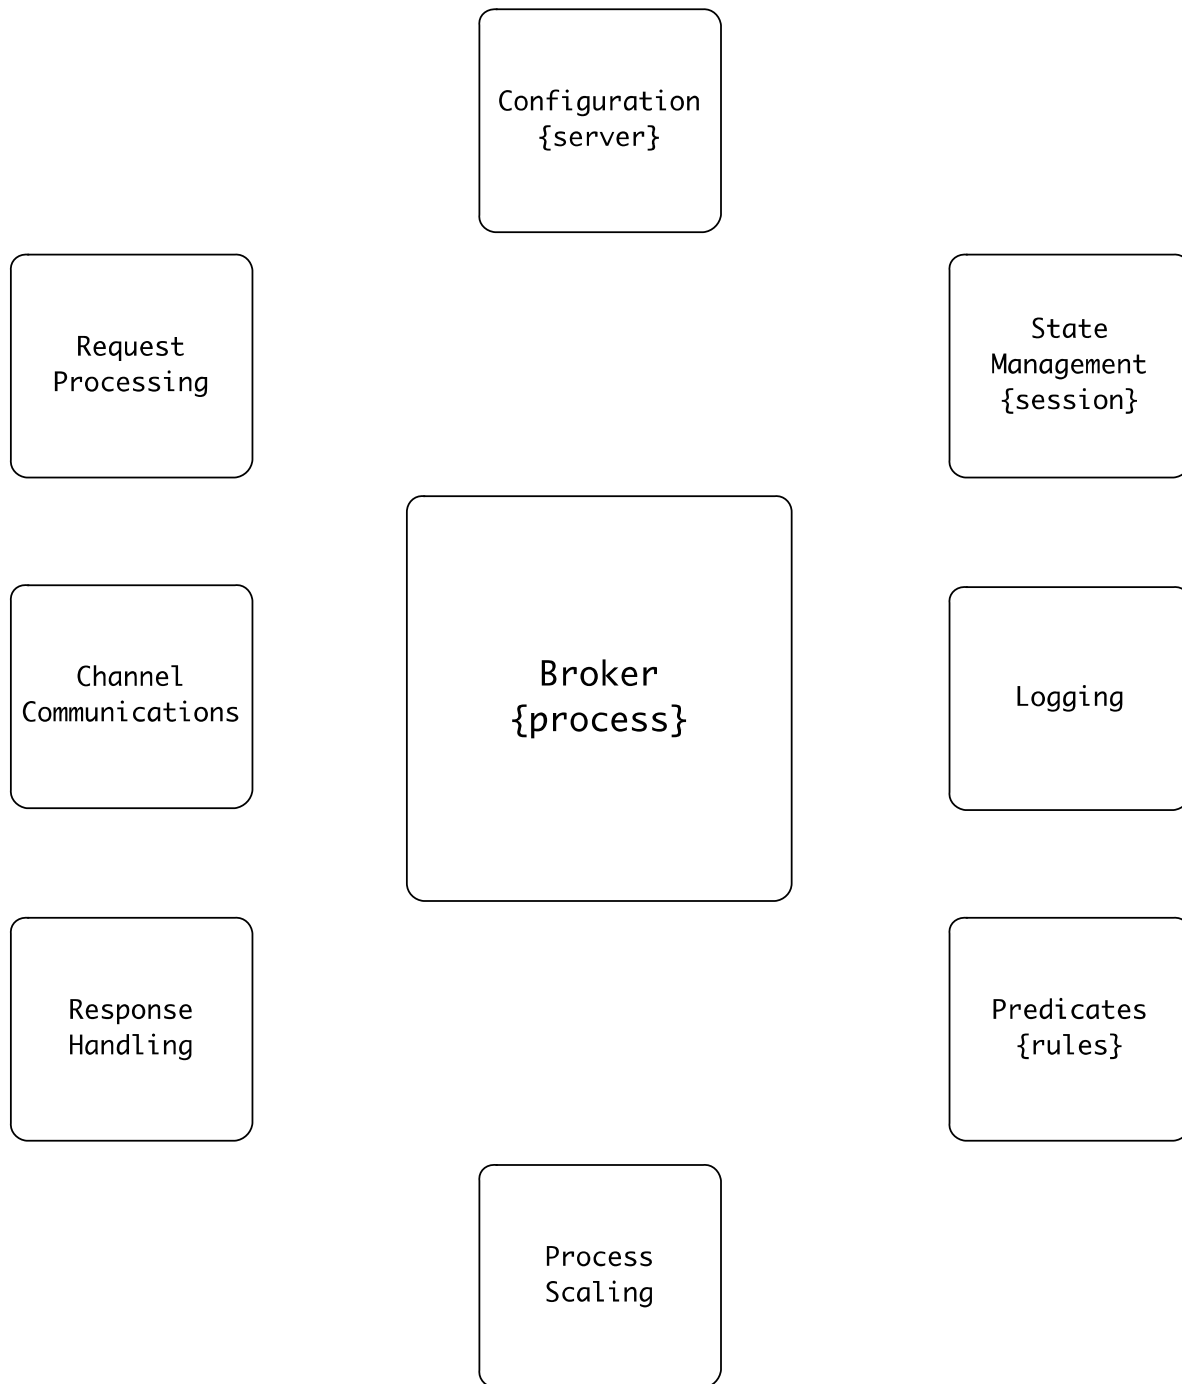
Value	Return Code Response	Description
0	0x00 Connection Accepted	Connection Accepted
1	0x01 Connection Refused	Unsupported MQTT Level
2	0x02 Connection Refused, Identifier Rejected	The Client Identifier is correct UTF-8 but not permitted
3	0x03 Connection Refused, Server Unavailable	The MQTT Service is unavailable
4	0x04 Connection Refused, Bad username or password	Data in the user name or password is malformed
5	0x05 Connection Refused, Not Authorized	Client not authorized to connect
6 - 255		Reserved for future use

Component Design

MQTT BROKER

Preliminaries

- Templates, templates, templates
- Policy-based Design
- Moderate use of Template Metaprogramming
- Implementation based upon First Principles
- Implemented as a daemon on Linux
- Personal research project exploring MQTT



IMPLEMENTATION DETAILS

Broker

```
template<typename HANDLE,  
        template <typename>  
        class Acceptor = Accept>  
  
class MqttBroker : public BasicDaemon  
                  <  
                  NetworkService  
                  <  
  
Acceptor<HANDLE>  
  
                  >  
                  >
```

BasicDaemon

```
template
<
    typename T,
    typename DaemonPolicy = daemon_policy,
    typename SigHandlerType = signals::SignalHandler,
    typename DaemonException = std::runtime_error
>
class BasicDaemon
```


Instantiating the Broker

```
// Read configuration etc...
```

```
MqttBroker<Handle> server;  
server.name(server_name);  
server.port(to_int(port));
```

```
// Launch the server, bootstrap and  
daemonize  
server.start();
```

Creating a Linux Daemon

- Fork a process
 - Call `fork()` - closes parent and makes child parent of init process
 - Call `setsid`, exit if it returns -1
- Call `fork()` again – a convenient shortcut
- Clear process' umask
- Change working directory to root directory
- Close all open file descriptors – careful with logging
- Use `dup2` to open descriptors 0, 1, and 2 to `/dev/null`

* *The Linux Programming Interface, Michael Kerrisk*

NetworkService

```
template<typename Accept,  
        typename Socket = SocketLifecycle<  
            typename
```

```
Accept::ServiceDescriptor>>
```

```
struct NetworkService : public Socket, Accept {
```

```
    void start() {  
        if (!is_listening()) {  
            this->connect();  
            this->listen();  
            this->accept();  
        }
```

```
    }
```

```
};
```

Waiting for Incoming Connections

```
// Back in the broker...
void _accept() override {

    while (true) {
        auto h = this->receive();
        if (this->is_good(h)) {

            if (!this->enable_read(h)) {
                log(str("Cannot continue: ") +
                    ErrorAdapter::get(errno));
            }
            // ... magic ...
        }
    }
}
```

Connection Received

```
auto buffer = handler.handle(h);
auto strategy =
    ResponseStrategy<char>::create(
        get_type(buffer), buffer);
auto exchange =
    mqtt::make_exchange(mqtt::make_message<char>
        (buffer.get()));

exchange.on_status_change(*this);
if (this->config_enabled(exchange)) {
    Task t(concurrent_work, exchange);
    _pool.submit(std::move(t));
}
```

What is an Exchange?

```
enum class ExchangeState {  
    CREATED, RUNNING, STOPPED,  
    BLOCKED, CANCELED, FINISHED  
};  
  
struct Exchange {  
    virtual void proceed() = 0;  
    virtual ExchangeState status() = 0;  
    virtual void status(ExchangeState) = 0;  
    virtual Configuration configuration() = 0;  
    virtual void configuration(Configuration) =  
0;  
};
```

An Exchange

- Task starts
- Exchange State:
 - Transitioned from CREATED to INPROGRESS
 - Event callback occurs on Broker
- Broker receives non-const reference to Exchange and can inspect for Client Id, enforce rules, and change state, e.g. block, stop or cancel
 - Depending on Exchange type, Broker may do one of several things
- Upon successful Exchange, final state is FINISHED
- Exchange communicates with Client and

Questions?