


Deploying C++ modules to 100s of millions of lines of code


Manuel Klimek (klimek@google.com)

CppCon 2016

Single TU compile time



Single TU compile time



The lay of the land

The Google Codebase

- multiple 100 MLOC C++ code
- ~same amount of generated code
- continuously integrated
- distributed and caching build system

Protocol Buffers

foo.proto:

```
message Foo {  
    optional int32 bar = 1;  
}
```

foo.pb.h:

```
class Foo {  
    int32 bar();  
    void set_bar(int32);  
    ...  
}
```

The build system

```
proto_library(name="p", srcs=["p.proto"])
```

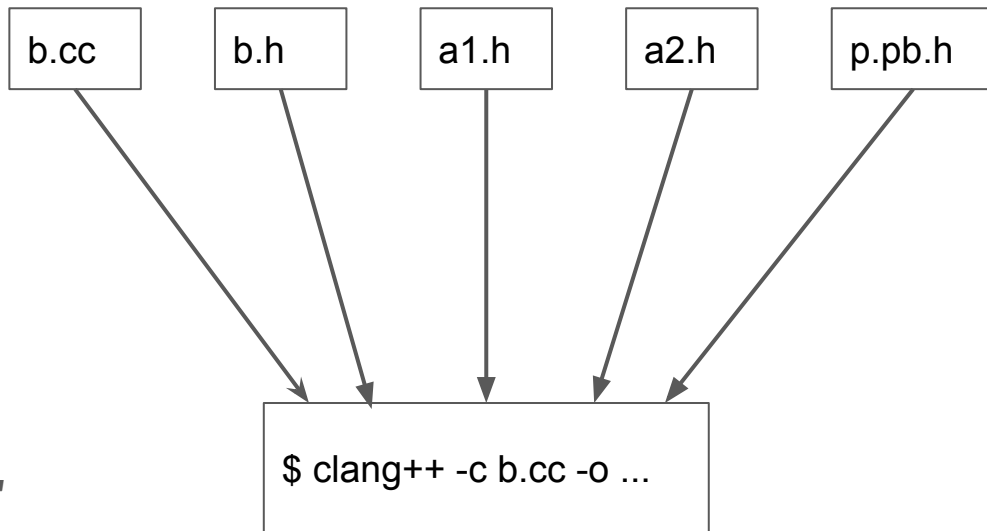
```
cc_library(name="a", hdrs=["a1.h", "a2.h"],  
           srcs=["a.cc"])
```

```
cc_library(name="b", hdrs=["b.h"], srcs=["b.cc"],  
           deps=["a", "p"])
```

b.cc:

```
#include "a1.h"  
#include "b.h"  
#include "other.h"  
#include "p.pb.h"
```


Send each file to distributed system to compile



b.cc:

```
#include "a1.h"  
#include "b.h"  
#include "other.h"  
#include "p.pb.h"
```

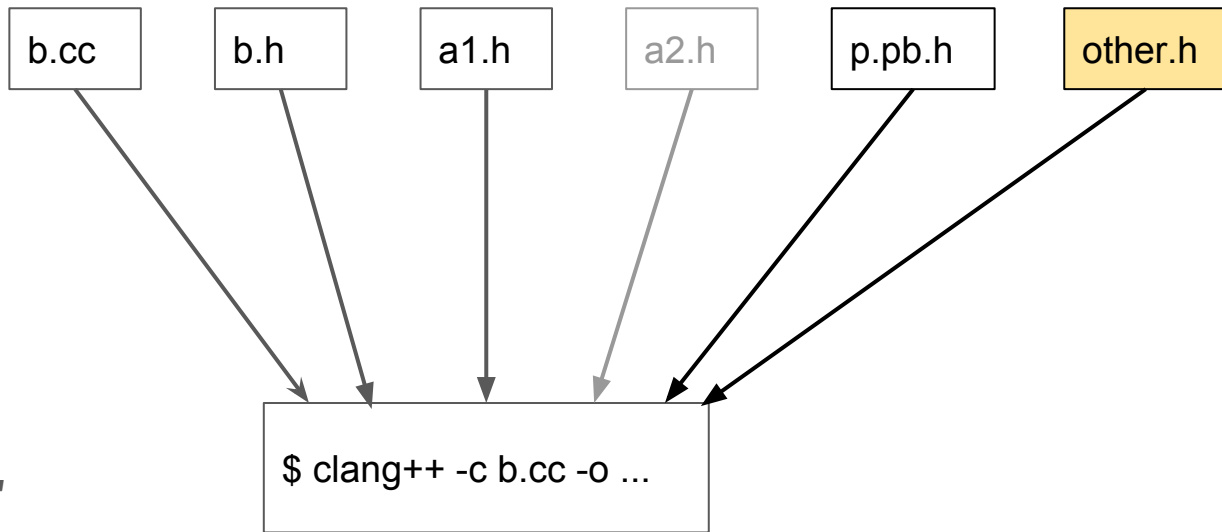

Scan code and prune unnecessary headers



b.cc:

```
#include "a1.h"  
#include "b.h"  
#include "other.h"  
#include "p.pb.h"
```

Scan code and add missing headers



b.cc:

```
#include "a1.h"  
#include "b.h"  
#include "other.h"  
#include "p.pb.h"
```

Modules? What modules?


- Clang's pre-TS modules implementation
- **no** changes to **syntax**
- changes to the **semantics**




Why not wait for modules and change all code?

- multiple 100 MLOC
- gain implementation insights
- back when we started, there was no modules TS proposal
- we had an acute problem

How bad did it get?



THE #1 PROGRAMMER EXCUSE
FOR LEGITIMATELY SLACKING OFF:
"MY CODE'S COMPILING."



What happened?

- was optimization getting slower? no
- was clang or gcc getting slower? no
- #lines of transitive include closure growing super-linearly!
(and the compilation scales mostly linearly with LOC)
- some TUs with > 10 MLOC transitive `#include` closure

Protocol buffers

a.proto:

```
message A {  
    optional string name = 1;  
}
```

a.pb.h

```
class A {  
    bool has_name() const;  
    void clear_name();  
    static const int kNameFieldNumber = 1;  
    const string& name() const;  
    void set_name(const string& v);  
    void set_name(const char* v);  
    void set_name(const char* v, size_t s);  
    string* mutable_name();  
    string* release_name();  
    void set_allocated_name(string* name);  
};  
  
inline bool A::has_name() const { ... }  
inline void A::clear_name() { ... }  
inline const string& A::name() const { ... }  
inline void A::set_name(const string& v) { ... }  
inline void A::set_name(const char* v) { ... }  
inline void A::set_name(const char* v, size_t s) { ... }  
inline string* A::mutable_name() { ... }  
inline string* A::release_name() { ... }  
inline void A::set_allocated_name(string* name) { ... }
```

Protocol buffers (solved!)

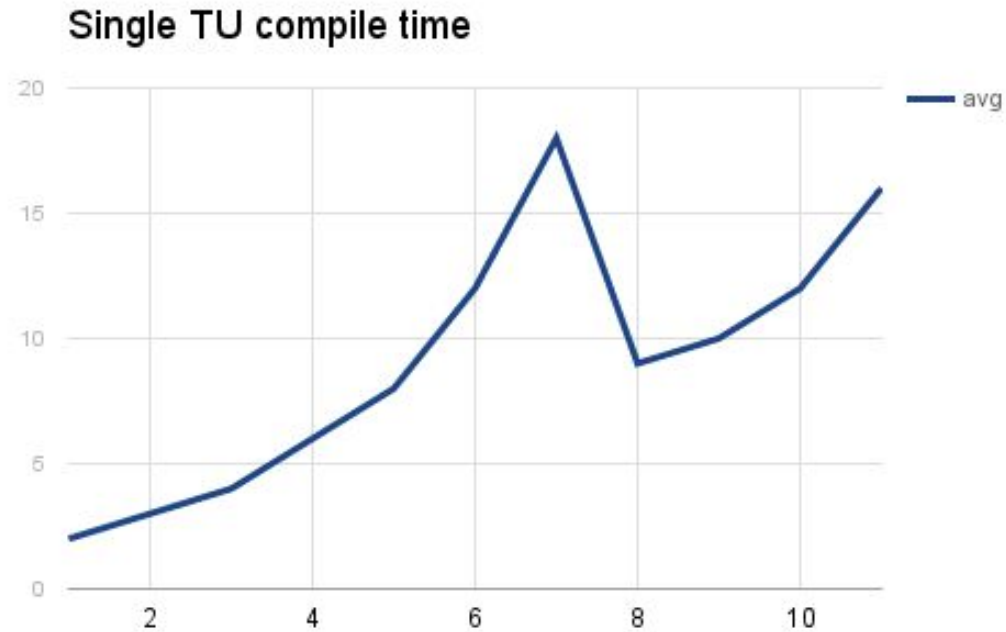
a.proto:

```
message A {  
    optional string name = 1;  
}
```

a.pb.h

```
class A {  
    bool has_name() const;  
    void clear_name();  
    static const int kNameFieldNumber = 1;  
    const string& name() const;  
    void set_name(const string& v);  
    void set_name(const char* v);  
    void set_name(const char* v, size_t s);  
    string* mutable_name();  
    string* release_name();  
    void set_allocated_name(string* name);  
};  
#ifndef NDEBUG  
inline bool A::has_name() const { ... }  
inline void A::clear_name() { ... }  
inline const string& A::name() const { ... }  
inline void A::set_name(const string& v) { ... }  
inline void A::set_name(const char* v) { ... }  
inline void A::set_name(const char* v, size_t s) { ... }  
inline string* A::mutable_name() { ... }  
inline string* A::release_name() { ... }  
inline void A::set_allocated_name(string* name) { ... }  
#endif /*NDEBUG*/
```

Solved! Solved?





Addressing the root cause (attempt 1)

- generated code needs to be **fast** (inline everything!)

b.pb.h:

```
#include "a.pb.h"

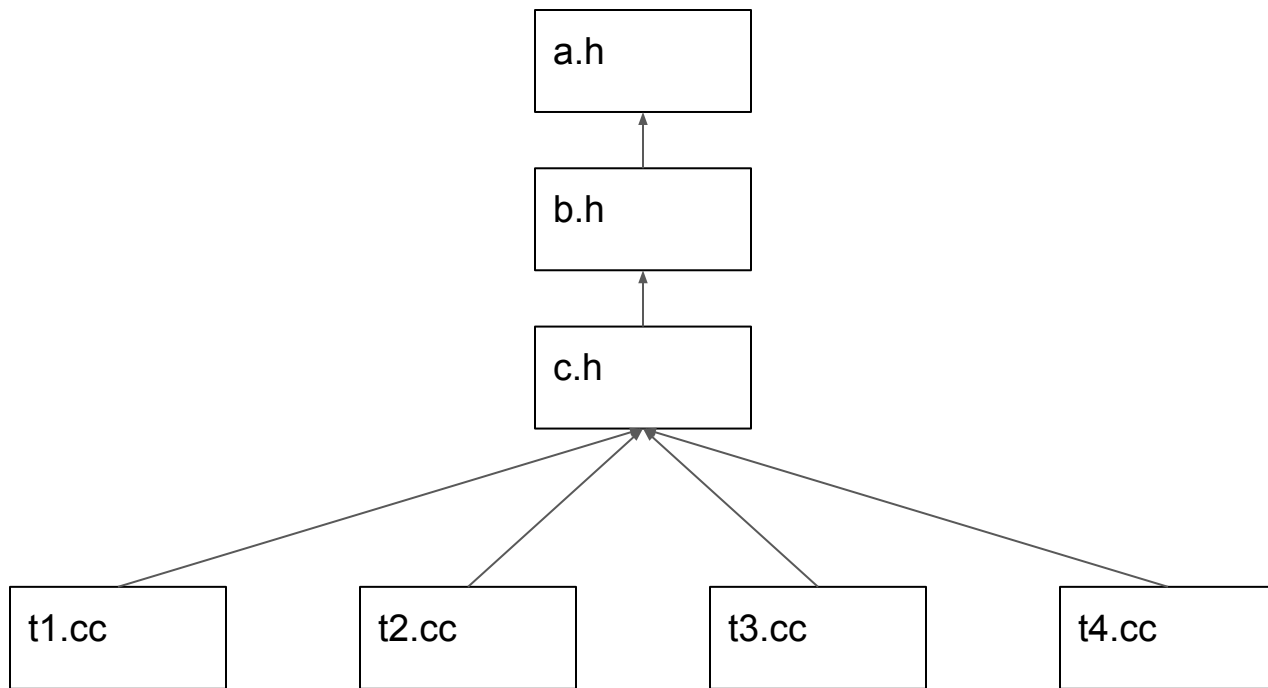
class B {
    ...
    A* mutable_a() {
        if (a_ == nullptr) a_ = new A();
        return a_;
    }
};
```

Addressing the root cause (attempt 2)

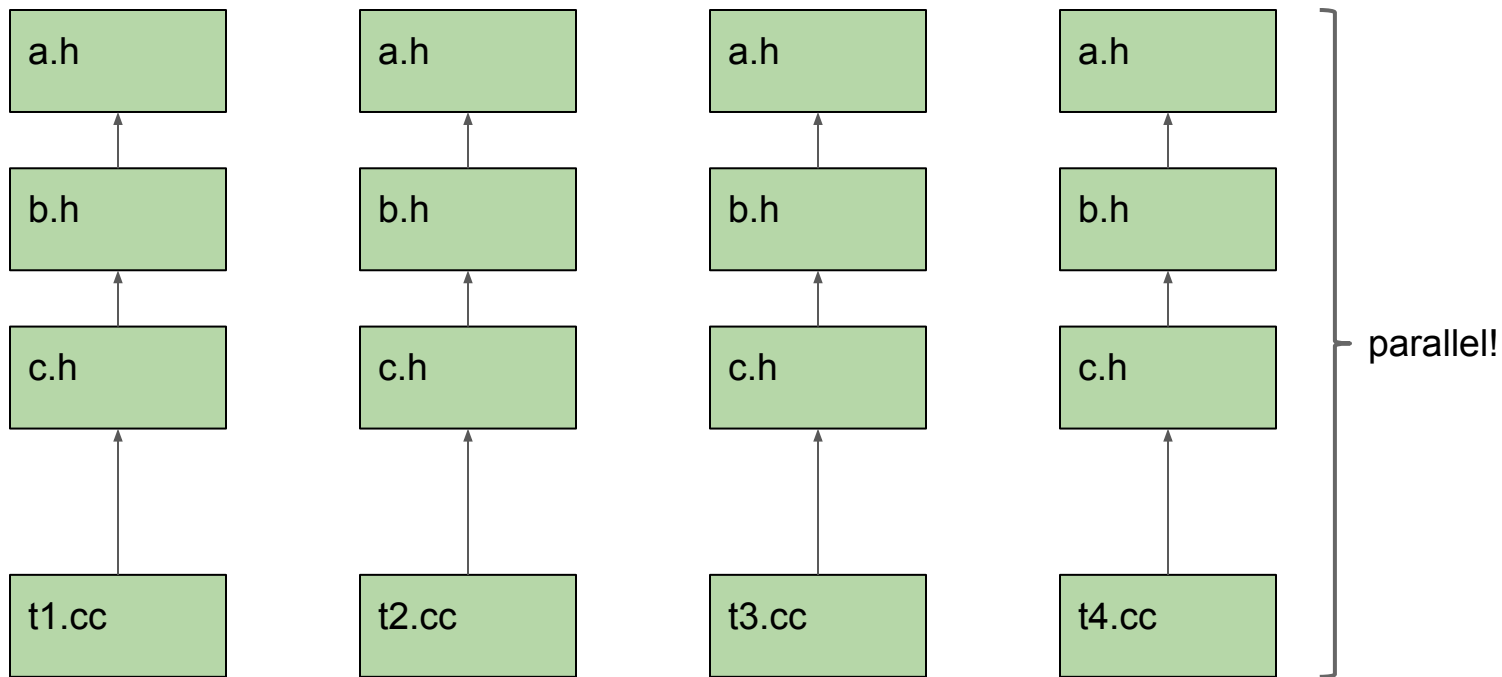
- need a compilation model that fundamentally doesn't scale superlinearly
- store AST and lazily load symbols
- > modules!

Let's predict what will happen...

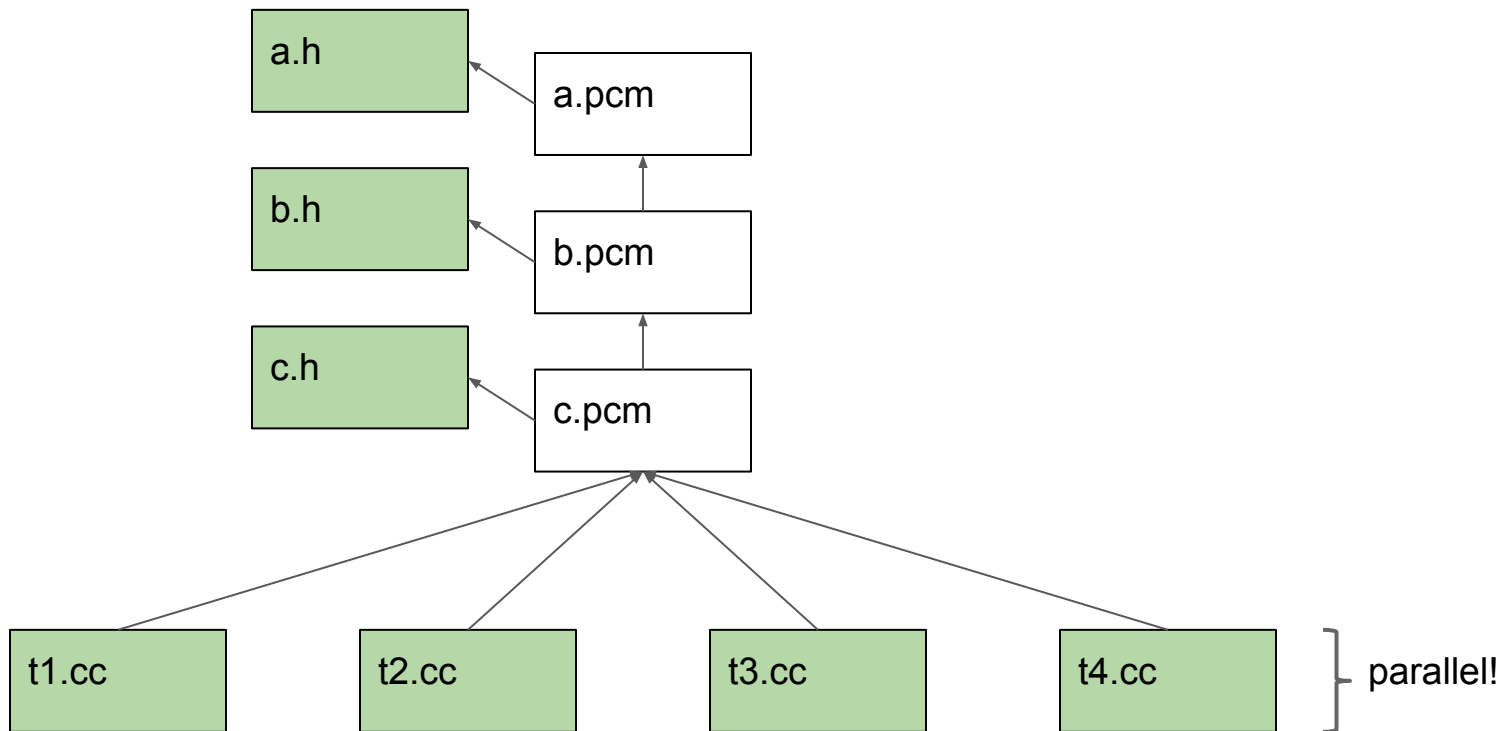
Exciting: a shiny new compilation model for C++




Exciting: a shiny new compilation model for C++




Exciting: a shiny new compilation model for C++



Exciting: a shiny new compilation model for C++



Exciting: a shiny new compilation model for C++



Exciting: a shiny new compilation model for C++

- > Longer critical path
- > Speedup of incremental compiles
- > Less CPU use overall



Selecting a subproblem

- protocol buffers:
 - generated code, controlled environment
 - largest problem

Describing a module to clang

```
proto_library(name="b", srcs=["b.proto"], deps=["a"])
```

```
module "b" {  
    export *  
    header "b.pb.h"  
    use "a"  
    use "stl"  
}
```


Non-modular headers

```
module "b" {  
  export *  
  header "b.pb.h" -> #include "message.h"  
  use "message"  
  use "stl"  
}
```

```
module "message" {  
  textual header "message.h"  
}
```

Semantic changes

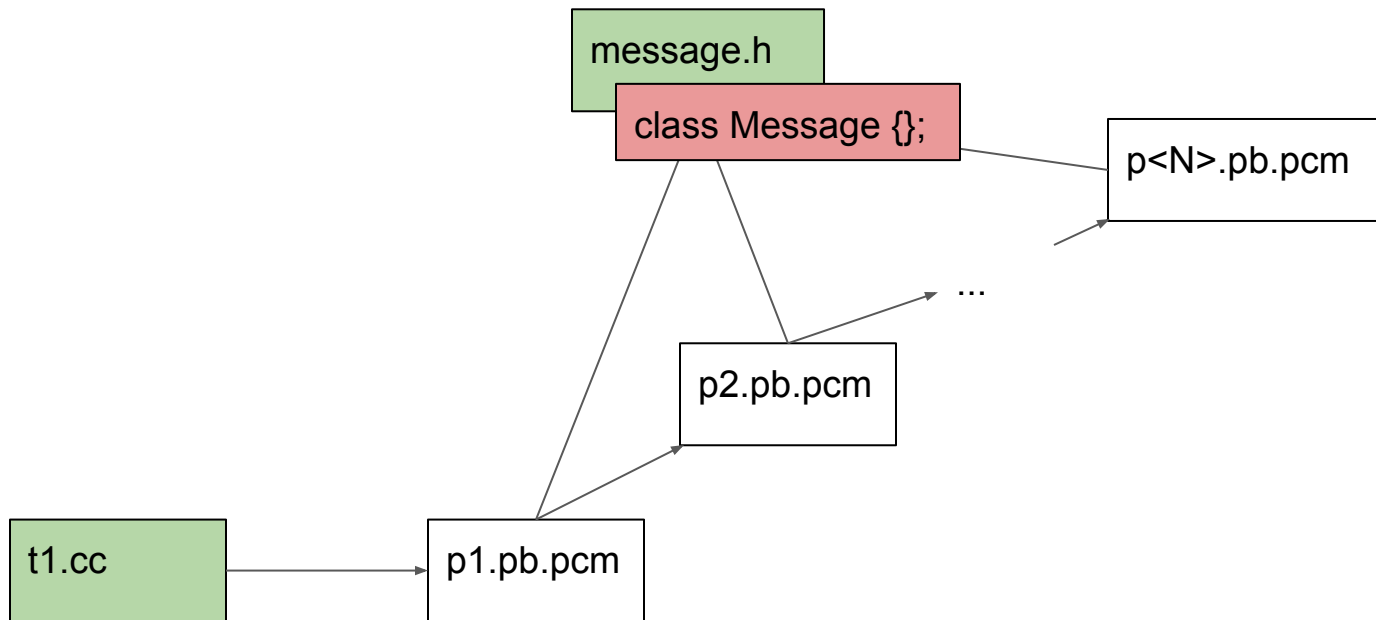
- headers must compile on their own
- more & different ODR violations diagnosed
- > headers must mean the same everywhere they're included from



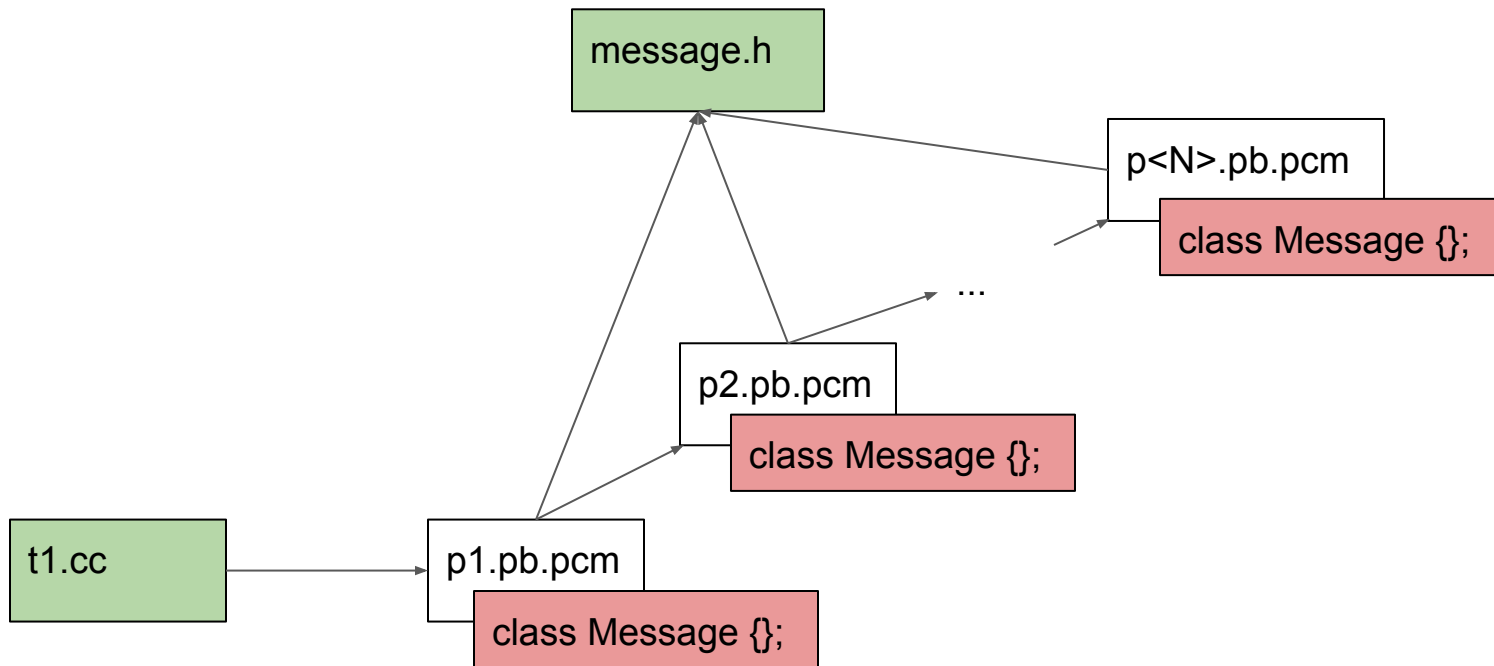
THINGS THAT BREAK

Things that break: modules twice as slow!

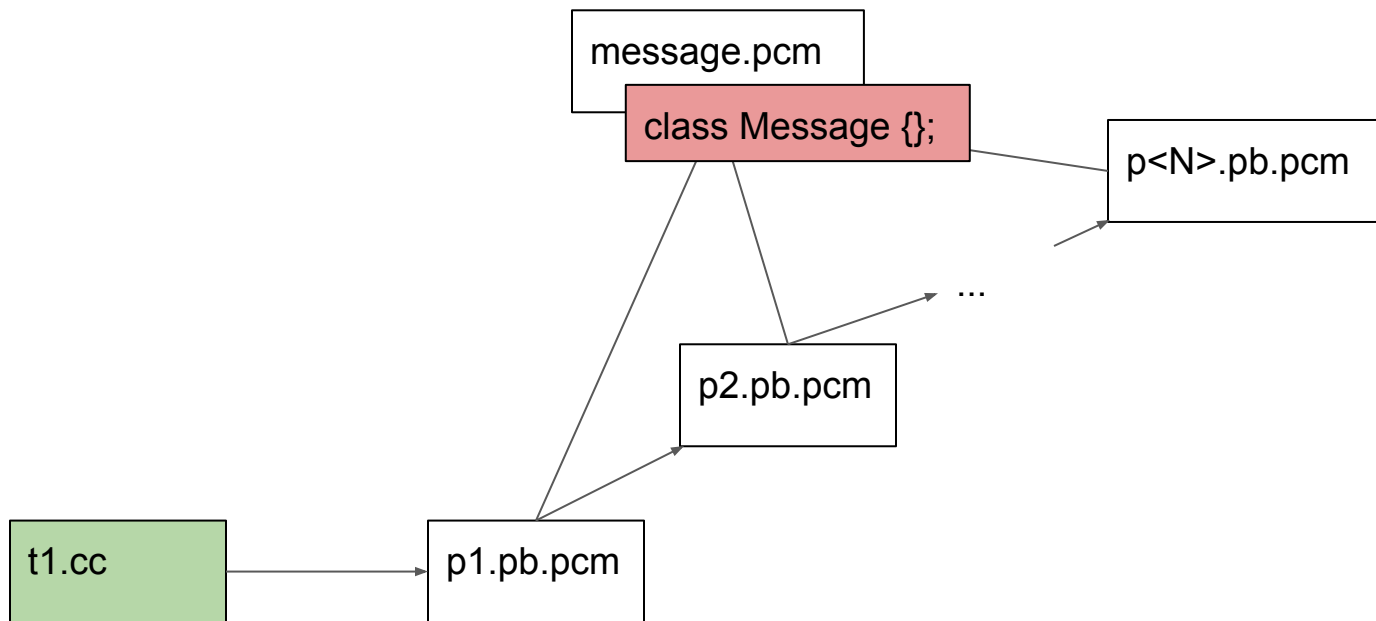
Things that break: modules twice as slow!



Things that break: modules twice as slow!



Things that break: modules twice as slow!



Problem: non-modular headers

- > performance problems when depending on many proto libraries (merging)
 - > each module contains a copy of the standard library, ...
- > some need for bottom-up modularization

Things that break: relying on #include order...

a.h:

```
int a();
```

b.h:

```
inline int b() { return a(); }
```

c.cc:

```
#include "a.h"
```

```
#include "b.h"
```

```
int main() { return b(); }
```

Things that break: _impl.h files

a.h:

```
struct C { int a(); };
```

...

```
#include "a_impl.h"
```

a_impl.h:

```
inline int C::a() { return 42; }
```

Things that break: C header

a.cc:

```
extern "C" {  
#include "b.h"  
}
```

Things that break: `__MODULE__`

- `__MODULE__` used by Arm C++ compiler conflicts with Clang's definition

Things that break: command line size

- explicit module files on command line cause command lines above the limits
- > pass only "top-level" module files
- > the calculation of those files makes the build system slower

Things that break: mixing configuration macros

a.h:

```
#include <cassert>

inline int a(int i) {
    assert(i > 0);
    return i+2;
}
```

b.cc:

```
#define NDEBUG
#include "a.h"

int main() { return a(-42); }
```


Configuration management

- need different module for each configuration
- theoretically: exponential explosion; in reality, a couple of configurations
- switch off modules when requiring different configurations
 - > automatically ignore modules


Flipping the switch

- headers are self-contained (> 10k changes)
- configuration management is under control
- distributed build system supports modules
- flip the switch!


Single TU compile time



Single TU compile time



Single TU compile time




:)

+ incremental rebuild better by up to 2x for the 99%ile

+ incremental rebuild average 10% better




The Big Regression




:(

- compiling a lot more translation units
- overall load increased due to increase in #compiles

Building too much



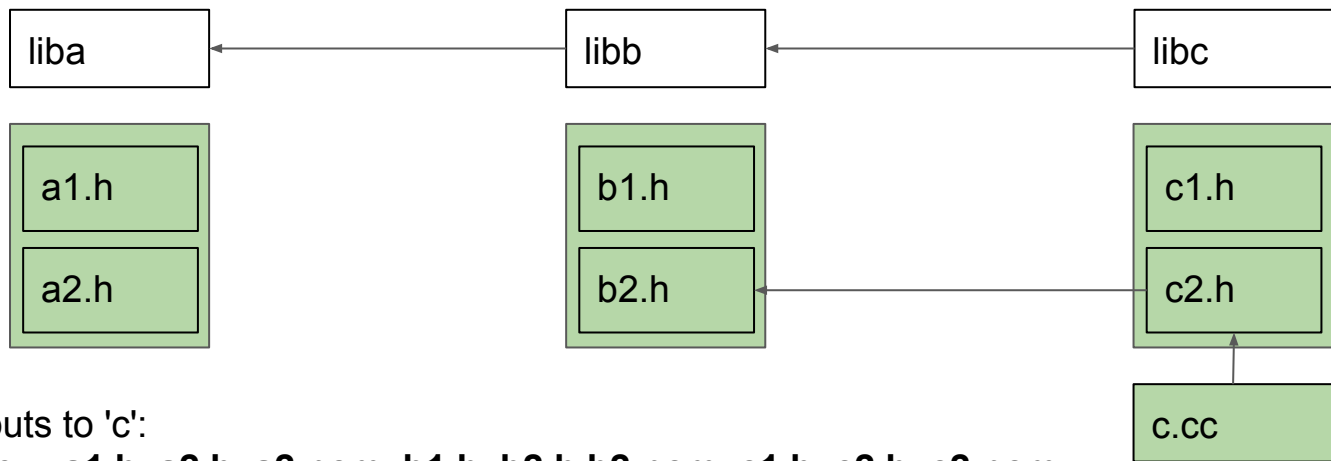
Building too much



Inputs to 'c':

- previously: b2.h, c2.h


Building too much



Inputs to 'c':

- now: **a1.h, a2.h, a2.pcm, b1.h, b2.h b2.pcm, c1.h, c2.h, c2.pcm**

Building too much




Inputs to 'c':

- now: **a1.h, a2.h, a2.pcm, b1.h, b2.h b2.pcm, c1.h, c2.h, c2.pcm**

- compile time **faster**, setup time **slower**

- lots more recompiles

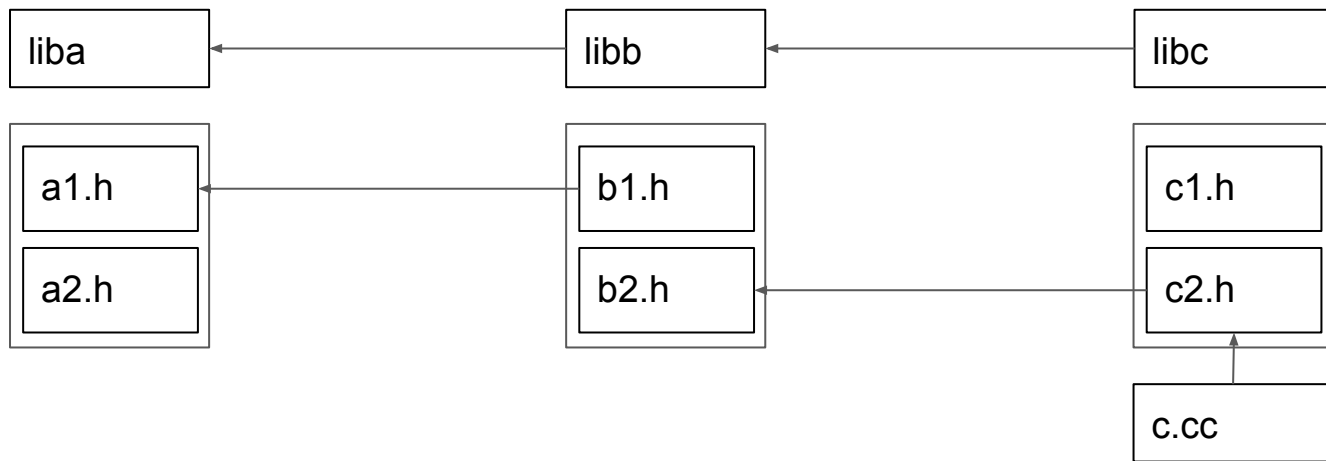
Building too much




Inputs to 'c':

- with include scanning: **b1.h, b2.h, b2.pcm**


Building too much




Building too much



Building too much



Building too much



Wrapping it up...

Results

- incremental compiles improved
 - up to 50% in the 99%-ile
 - avg ~10%
- clean build times ~equivalent
- overall build load increased due to increasing #compiles
- unlocked loads of optimization potential
- prepares the code base for C++ modules

What worked

- we were lucky: style guide -> most things pretty good
- tools to fix things at scale -> fix broken code
- build system models headers -> build system integration doable

What was problematic


- performance improvement not as large as expected
- broken code sometimes really hard to fix
- lots of effort tweaking the distributed build system
- bleeding edge of clang's features -> bugs

Questions?!

Rolling out Modules

1. Implement modules support in the build system
2. Make headers self-contained (-x c++-header)
3. Get configurations under control
4. Switch on modules
5. Fix things that break
6. Optimize your build system

Growing linker input



The build system

```
proto_library(name="p", srcs=["p.proto"])
```

```
cc_library(name="a", hdrs=["a1.h", "a2.h"],  
           srcs=["a.cc"])
```

```
cc_library(name="b", hdrs=["b.h"], srcs=["b.cc"],  
           deps=["a", "p"])
```

b.cc:

```
#include "a1.h"  
#include "b.h"  
#include "other.h"  
#include "p.pb.h"
```

