



# I Just Wanted a Random Integer!

Cppcon 2016

Cheinan Marks

Spiral Genetics, Inc.



I Just Wanted a Random Integer!

ATTCTGTAGCGTGCATGCATGTTCGAT



# I Just Wanted a Random Integer!

ATTCTGTAGCGTGCATGCATGTCTGAT

- Needed random sized DNA reads
- Wanted good coverage of size range
- Wanted to test edge cases



# I Just Wanted a Random Integer!

```
#include <cstdlib>
```

```
auto r = std::rand() % 100;
```



# I Just Wanted a Random Integer!

```
#include <cstdlib>  
auto r = std::rand() % 100;
```

```
man -s3 rand
```

## NOTES

The versions of `rand()` and `srand()` in the Linux C Library use the same random number generator as `random(3)` and `srandom(3)`, so the lower-order bits should be as random as the higher-order bits.

However, on older `rand()` implementations, and on current implementations on different systems, the lower-order bits are much less random than the higher-order bits. Do not use this function in applications intended to be portable when good randomness is needed. (Use `random(3)` instead.)



# I Just Wanted a Random Integer!

```
#include <cstdlib>  
auto r = std::rand() % 100;
```

```
man -s3 rand
```

## NOTES

The versions of `rand()` and `srand()` in the Linux C Library use the same random number generator as `random(3)` and `srandom(3)`, so the lower-order bits should be as random as the higher-order bits.

However, on older `rand()` implementations, and on current implementations on different systems, **the lower-order bits are much less random than the higher-order bits. Do not use this function in applications intended to be portable when good randomness is needed. (Use `random(3)` instead.)**



# I Just Wanted a Random Integer!

- It's just a unit test
- `std::rand` should be good enough
- `%` should be good enough
- `<random>` is too complicated



# I Just Wanted a Random Integer!

- It's just a unit test
- ~~std::rand~~ should be good enough **NOPE**
- ~~%~~ should be good enough **NOPE**
- <random> is too complicated **Err, maybe**





# I Just Wanted a Random Integer!

- It's just a unit test
- ~~std::rand~~ should be good enough **NOPE**
- ~~%~~ should be good enough **NOPE**
- ~~<random>~~ is too complicated **Err, maybe**
- If you want even coverage of edge cases
- Then you **must** do better



# I Just Wanted a Random Integer!

1. Watch STL's 2013 Going Native Talk
2. Use `<random>` and Mersenne Twister
3. Seed with `std::random_device` (entropy)
4. Use `std::uniform_int_distribution`
5. Profit



# I Just Wanted a Random Integer!

## Questions

- Entropy – what is it?
- What is `std::random_device`?
- Why avoid the stack for `std::mt19937`?
- Is `std::uniform_int_distribution` cheap to construct and use?



# I Just Wanted a Random Integer!

- Entropy?
- Comes from `std::random_device`
- Blocks when it runs out
- Wait, what?!?



# I Just Wanted a Random Integer!

- Entropy is randomness
- Computers are not by nature random
- Computers can be pseudorandom
- Entropy is “true randomness”
- Hard to generate on a computer



# I Just Wanted a Random Integer!

- Entropy?
- Comes from `std::random_device`
- Blocks when it runs out
- Wait, what?!?



# I Just Wanted a Random Integer!

- Entropy?
- Comes from `std::random_device`
- Blocks when it runs out
- Shrink person with die and whiteboard
- Nanoperson rolls die
- Nanoperson writes result on whiteboard
- Nano person returns number on demand



# I Just Wanted a Random Integer!

- It takes finite time to roll die
- If demand > supply, entropy runs out
- Entropy generator might block until more entropy is available





# I Just Wanted a Random Integer!

- `std::random_device`
- Can be used as a temporary
- Can use it as a generator
- `operator()` returns  $\text{min}() \leq \text{RN} \leq \text{max}()$
- Hardware and implementation dependent
- CAUTION: Can throw
- Might be slow, might be pseudorandom
- Uses `/dev/urandom` on my machine



# I Just Wanted a Random Integer!

- Generators
- Use `std::mt19937` or `std::mt19937_64`
- Seed with `std::random_device`
- Pseudorandom
- Deterministic
- Fast



# I Just Wanted a Random Integer!

## Generate billion random ints

<code>std::random_device</code>	<code>std::mt19937</code>
44.3 seconds	3.6 seconds



# I Just Wanted a Random Integer!

## Generate billion random ints

<code>std::random_device</code>	<code>std::mt19937</code>
44.3 seconds	3.6 seconds

- `std::random_device` may be hardware
- Multithreading behavior unclear
- `std::mt19937` can be thread local
- Initialize correctly!



# I Just Wanted a Random Integer!

1. Watch STL's 2013 Going Native Talk
2. Use Mersenne Twister in `<random>`
3. Seed with `std::random_device` (entropy)
4. Use `std::uniform_distribution`
5. Profit Think



# I Just Wanted a Random Integer!

`std::mt19937` is good, but...



# I Just Wanted a Random Integer!

`std::mt19937` is good, but...

5000 bytes on the stack

Slow to initialize



# I Just Wanted a Random Integer!

`std::mt19937` is good, but...

5000 bytes on the stack

Slow to initialize

14.991 seconds for 1,000,000 inits

Slower than `std::random_device`!





# I Just Wanted a Random Integer!

`std::mt19937` is good, but...

5000 (2504) bytes on the stack

Slow to initialize

15.0 seconds for 1,000,000 inits

```
void f() {  
    std::mt19937 g(std::random_device{}());  
    auto rn = g();  
}
```



# I Just Wanted a Random Integer!

std::mt19937 is good, but...

5000 (2504) bytes on the stack

Slow to initialize

15.0 seconds for 1,000,000 inits

```
void f() {  
    static std::mt19937 g(std::random_device{}());  
    auto rn = g();  
}
```



# I Just Wanted a Random Integer!

## Guidelines Make Sense

- Use `std::random_device` to seed
- Keep `std::mt19937` off stack



# I Just Wanted a Random Integer!

## Guidelines Make Sense

- Use `std::random_device` to seed
- Keep `std::mt19937` off stack
- It can be static, thread local
- On stack, but beware construction



# I Just Wanted a Random Integer!

## Guidelines Make Sense

- Use `std::random_device` to seed
- Keep `std::mt19937` off stack
- It can be static, thread local
- On stack, but beware construction
- `std::minstd_rand` is much faster
- `std::minstd_rand` cycle is smaller



# I Just Wanted a Random Integer!

## Guidelines Make Sense

- Use `std::random_device` to seed
- Keep `std::mt19937` off stack
- It can be static, thread local
- On stack, but beware construction
- `std::minstd_rand` is much faster
- `std::minstd_rand` cycle is smaller

random_device	mt19937	minstd_rand
44.3 seconds	3.6 seconds	4.7 seconds



# I Just Wanted a Random Integer!

## Guidelines Make Sense

- Use `std::random_device` to seed
- Keep `std::mt19937` off stack
- It can be static, thread local
- On stack, but beware construction
- ~~`std::minstd_rand` is much faster~~
- `std::minstd_rand` cycle is smaller

random_device	mt19937	minstd_rand
44.3 seconds	3.6 seconds	4.7 seconds



# I Just Wanted a Random Integer!

Marin Mersenne 1588-1648





# I Just Wanted a Random Integer!

Marin Mersenne 1588-1648

$$M_n = 2^n - 1, \text{ prime } n$$

$$M_3 = 7, M_7 = 127$$

...

$$M_{74,207,281} \sim 10^{22,338,618}$$



# I Just Wanted a Random Integer!

Marin Mersenne 1588-1648

$$M_n = 2^n - 1, \text{ prime } n$$

$$M_3 = 7, M_7 = 127$$

...

$$M_{74,207,281} \sim 10^{22,338,618}$$

Why `std::mt19937`?

$$\text{Period} = M_{19937} \sim 10^{6002}$$





# I Just Wanted a Random Integer!

Marin Mersenne 1588-1648

$$M_n = 2^n - 1, \text{ prime } n$$

$$M_3 = 7, M_7 = 127$$

...

$$M_{74,207,281} \sim 10^{22,338,618}$$

Why std::mt19937?

$$\text{Period} = M_{19937} \sim 10^{6002}$$



```
typedef mersenne_twister_engine<uint_fast32_t, 32, 624, 397, 31,  
0x9908b0dfUL, 11, 0xffffffffUL, 7, 0x9d2c5680UL, 15,  
0xefc60000UL, 18, 1812433253UL> mt19937;
```

# I Just Wanted a Random Integer!

Marin Mersenne 1588-1648

$$M_n = 2^n - 1, \text{ prime } n$$

$$M_3 = 7, M_7 = 127$$

...

$$M_{74,207,281} \sim 10^{22,338,618}$$

Why `std::mt19937`?

$$\text{Period} = M_{19937} \sim 10^{6002}$$



M. Matsumoto and T. Nishimura, Mersenne Twister: A 623-Dimensionally Equidistributed Uniform Pseudo-Random Number Generator, ACM Transactions on Modeling and Computer Simulation, Vol. 8, No. 1, January 1998, pp 3-30.



# I Just Wanted a Random Integer!

## Guidelines Make Sense

- Use `std::random_device` to seed
- Keep `std::mt19937` off stack
- It can be static, thread local
- On stack, but beware construction
- ~~`std::minstd_rand` is much faster~~
- `std::minstd_rand` cycle is smaller

random_device	mt19937	minstd_rand
44.3 seconds	3.6 seconds	4.7 seconds



# I Just Wanted a Random Integer!

Guidelines Make Sense

- Not so fast...



I Just Wanted a Random Integer!

STL Inspiration





# I Just Wanted a Random Integer!

STL Inspiration

Melissa O'Neill, Harvey Mudd College

PCG <http://www.pcg-random.org>

[https://www.reddit.com/r/programming/comments/2momvr/pcg\\_a\\_family\\_of\\_better\\_random\\_number\\_generators/](https://www.reddit.com/r/programming/comments/2momvr/pcg_a_family_of_better_random_number_generators/)

Search reddit.com for “PCG random”

<http://preview.tinyurl.com/hyfb73l>





# I Just Wanted a Random Integer!

Melissa O'Neill, PCG

- Smaller (16 bytes) than `std::mt19937`
- C++ `<random>` compatible library
- <https://github.com/imneme/pcg-cpp>
- Faster than `std::mt19937`



# I Just Wanted a Random Integer!

Melissa O'Neill, PCG

- Smaller (16 bytes) than `std::mt19937`
- C++ `<random>` compatible library
- <https://github.com/imneme/pcg-cpp>
- Faster than `std::mt19937`

random_device	mt19937	pcg_32
44.3 seconds	3.6 seconds	1.5 seconds

- Not in the standard
- Much less real-life experience



# I Just Wanted a Random Integer!

`std::uniform_int_distribution`



# I Just Wanted a Random Integer!

`std::uniform_int_distribution`

Instantiate in inner loop?



# I Just Wanted a Random Integer!

`std::uniform_int_distribution`

Instantiate in inner loop?

Look into implementation

Stores two template arguments

Constructor is trivial



# I Just Wanted a Random Integer!

`std::uniform_int_distribution`

Instantiate in inner loop?

Look into implementation

Stores two template arguments

Constructor is trivial

Operator() branches and calculates

Not much code



# I Just Wanted a Random Integer!

```
static std::random_device entropySource;  
static std::mt19937 randGenerator(entropySource());  
std::uniform_int_distribution<int> theIntDist(0, 99);  
  
for (auto i = 0; i < 1'000'000'000; i++) {  
    volatile auto r = theIntDist(randGenerator);  
}  
  
// 23.4 seconds
```



# I Just Wanted a Random Integer!

```
static std::random_device entropySource;  
static std::mt19937 randGenerator(entropySource());  
  
for (auto i = 0; i < 1'000'000'000; i++) {  
    std::uniform_int_distribution<int> theIntDist(0, 99);  
    volatile auto r = theIntDist(randGenerator);  
}
```





# I Just Wanted a Random Integer!

```
static std::random_device entropySource;  
static std::mt19937 randGenerator(entropySource());  
  
for (auto i = 0; i < 1'000'000'000; i++) {  
    std::uniform_int_distribution<int> theIntDist(0, 99);  
    volatile auto r = theIntDist(randGenerator);  
}  
  
// 5.1 seconds
```



# I Just Wanted a Random Integer!

Constructor Outside Loop	Constructor Inside Loop
23.4 seconds	5.1 seconds



# I Just Wanted a Random Integer!

Constructor Outside Loop	Constructor Inside Loop
23.4 seconds	5.1 seconds
No Optimization	No Optimization
49.4 seconds	57.5 seconds



# I Just Wanted a Random Integer!

```
const __uctype __urange
= __uctype(__param.b()) - __uctype(__param.a());

__uctype __ret;

if (__urngrange > __urange)
{
    // downscaling
    const __uctype __uerange = __urange + 1; // __urange can be zero
    const __uctype __scaling = __urngrange / __uerange;
    const __uctype __past = __uerange * __scaling;
    do
        __ret = __uctype(__urng()) - __urngmin;
    while (__ret >= __past);
    __ret /= __scaling;
}
```



# I Just Wanted a Random Integer!

```
const __uctype __urange
= __uctype(__param.b()) - __uctype(__param.a());

__uctype __ret;

if (__urngrange > __urange)
{
    // downscaling
    const __uctype __uerange = __urange + 1; // __urange can be zero
    const __uctype __scaling = __urngrange / __uerange;
    const __uctype __past = __uerange * __scaling;
    do
        __ret = __uctype(__urng()) - __urngmin;
    while (__ret >= __past);
    __ret /= __scaling;
}
```



I Just Wanted a Random Integer!

Guidelines



# I Just Wanted a Random Integer!

## Guidelines

- Use your engineering judgment
- `<random>` is safe
- PCG is fast, small and simple
- Combine PCG with `<random>`
- Always measure. Always!



# I Just Wanted a Random Integer!

## Guidelines / Conclusions

- Use your engineering judgment
- `<random>` is safe
- PCG is fast, small and simple
- Combine PCG with `<random>`
- Always **measure**. Always!





# I Just Wanted a Random Integer!

## Conclusions

`std::random_device::operator()` to generate random numbers

`std::mt19937` or PCG generators

Distributions are cheap to construct

Distributions are cheap to use

C++17 has `std::sample`

Benchmark your code



I Just Wanted a Random Integer!

?